

# imdb\_movie

July 20, 2019

Author: Shashi Bala Table of contents

```
<ol>
  <li><a href="#about_dataset">About the dataset</a></li>
  <li><a href="#exploring_the_data">Data Exploration</a></li>
  <li><a href="#cleaning_the_data">Data Cleaning</a></li>
  <li><a href="#transforming_the_data">Data Transformation</a></li>
  <li><a href="#using_parameteric_and_non_parametric_models">Modeling</a></li>
  <li><a href="#recommendations">Recommendation</a></li>
</ol>
```

Import the Following Libraries:

```
<li> <b>numpy (as np)</b> </li>
<li> <b>pandas (as pd)</b> </li>
<li> <b>pandas_profiling</b> </li>
<li> <b>matplotlib.pyplot (as plt)</b> </li>
<li> <b>seaborn (as sns)</b> </li>
<li> <b>warnings</b> </li>
<li> <b>os</b> </li>
<li> <b>series, DataFrame</b> from <b>pandas</b> </li>
<li> <b>stats</b> from <b>scipy.stats</b> </li>
<li> <b>train_test_split</b> from <b>sklearn.model_selection</b> </li>
<li> <b>LinearRegression</b> from <b>sklearn.linear_model</b> </li>
<li> <b>r2_score</b> from <b>sklearn.metrics</b> </li>
<li> <b>statsmodels.api (as sm)</b> </li>
<li> <b>KNeighborsRegressor</b> from <b>sklearn.neighbors</b> </li>
<li> <b>mean_squared_error</b> from <b>sklearn.metrics</b> </li>
<li> <b>neighbors</b> from <b>sklearn</b> </li>
<li> <b>sqrt</b> from <b>math</b> </li>
<li> <b>RandomForestRegressor</b> from <b>sklearn.ensemble</b> </li>
```

```
In [2]: import pandas as pd
import pandas_profiling
from pandas import Series, DataFrame
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
```

```

import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import LabelEncoder
import scipy.stats as stats
import os
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import statsmodels.api as sm
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error
from math import sqrt
from sklearn.ensemble import RandomForestRegressor

```

<h2>1. About the dataset</h2>

The dataset is imdb\_data\_v2 which is in csv file. It contains 38 variables for 5787 movies, sp

<h2>2. Data Exploration</h2>

### 2.1. Data Load

```
In [3]: print(os.listdir("C:\Data Science_Interview\Dataset"))
```

```
['Data Scientist imdb_data_v2.csv']
```

We read in the data we've saved, passing the column names

```
In [4]: df = pd.read_csv('C:\Data Science_Interview\Dataset\Data Scientist imdb_data_v2.csv')
```

```
In [5]: print('Training data shape: ', df.shape)
```

```
Training data shape: (5787, 38)
```

Let's check out the first few rows of data

```
In [6]: df.head()
```

```
Out[6]:
```

	id	stock_market_idx	days_since_last_tweet	pre_screen_viewers	\
0	1	1102	67	18	
1	2	1117	70	18	
2	3	1000	90	11	
3	4	1007	35	10	
4	5	1128	85	20	

	characters_per_longest_review	priority	\
0	1181	4	
1	1196	4	
2	1125	4	

3		1127		4
4		1072		4

	longest_facebook_comment_review_char	color	director_name	\
0	250	Color	James Cameron	
1	740	Color	Gore Verbinski	
2	1779	Color	Sam Mendes	
3	1074	Color	Christopher Nolan	
4	813	NaN	Doug Walker	

	num_critic_for_reviews	...	country	content_rating	website_score	\
0	723.0	...	USA	PG-13	7.9	
1	302.0	...	USA	PG-13	7.1	
2	602.0	...	UK	PG-13	6.8	
3	813.0	...	USA	PG-13	8.5	
4	NaN	...	NaN	NaN	7.1	

	budget	weighted_budget	title_year	actor_2_facebook_likes	\
0	237000000.0	236999000	2009.0	936.0	
1	300000000.0	299999000	2007.0	5000.0	
2	245000000.0	244999000	2015.0	393.0	
3	250000000.0	249999000	2012.0	23000.0	
4	NaN	-1000	NaN	12.0	

	aspect_ratio	movie_facebook_likes	imdb_score
0	1.78	33000	7.9
1	2.35	0	7.1
2	2.35	85000	6.8
3	2.35	164000	8.5
4	NaN	0	7.1

[5 rows x 38 columns]

## 0.0.1 Data Profiling

```
In [7]: df.profile_report()
```

```
<IPython.lib.display.IFrame at 0x1fa28223cf8>
```

Out [7]:

We have 5787 observations of 38 variables in which 21 variables are numeric and 11 variables are categorical. The response variable “imdb\_score” is numerical, and the predictors are mixed with numerical and categorical variables.

### 2.2. Remove Duplicates

In the IMDB dataset, There is 744 (12.9%) duplicate rows. I want to remove the 744 duplicated rows and keep the unique ones.

```
In [8]: df1 = df
```

```
In [9]: #drop the duplicates  
df1.drop_duplicates(inplace=True)  
# Check if done  
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 5043 entries, 0 to 5042  
Data columns (total 38 columns):  
id                    5043 non-null int64  
stock_market_idx      5043 non-null int64  
days_since_last_tweet 5043 non-null int64  
pre_screen_viewers    5043 non-null int64  
characters_per_longest_review 5043 non-null int64  
priority              5043 non-null int64  
longest_facebook_comment_review_char 5043 non-null int64  
color                 5024 non-null object  
director_name         4939 non-null object  
num_critic_for_reviews 4993 non-null float64  
duration              5028 non-null float64  
director_facebook_likes 4939 non-null float64  
actor_3_facebook_likes 5020 non-null float64  
actor_2_name          5030 non-null object  
actor_1_facebook_likes 5036 non-null float64  
gross                 4159 non-null float64  
genres                5043 non-null object  
actor_1_name          5036 non-null object  
movie_title           5043 non-null object  
num_voted_users       5043 non-null int64  
cast_total_facebook_likes 5043 non-null int64  
made_up_column        5043 non-null float64  
actor_3_name          5020 non-null object  
facenumber_in_poster  5030 non-null float64  
plot_keywords         4890 non-null object  
movie_imdb_link        5043 non-null object  
num_user_for_reviews  5022 non-null float64  
language              5031 non-null object  
country               5038 non-null object  
content_rating         4740 non-null object  
website_score          5043 non-null float64  
budget                4551 non-null float64  
weighted_budget       5043 non-null int64  
title_year            4935 non-null float64  
actor_2_facebook_likes 5030 non-null float64  
aspect_ratio          4714 non-null float64  
movie_facebook_likes   5043 non-null int64  
imdb_score             5043 non-null float64
```

```
dtypes: float64(15), int64(11), object(12)
memory usage: 1.5+ MB
```

## <h2>3. Data Cleaning</h2>

### 3.1 Missing Values

We can quickly check if we have any null values in our data

```
In [10]: def mis_values(df1):
        mis_value = df1.isnull().sum()
        mis_value_per = 100 * df1.isnull().sum() / len(df1)
        mis_value_column = pd.concat([mis_value, mis_value_per], axis=1)
        mis_val_tab_rename_cols = mis_value_column.rename(columns = {0 : 'Missing Values'})
        mis_val_tab_rename_cols = mis_val_tab_rename_cols[mis_val_tab_rename_cols.iloc[:,0] != 0]
        print ("Your selected dataframe has " + str(df.shape[1]) + " columns.\n"

                "There are " + str(mis_val_tab_rename_cols.shape[0]) +
                " cols that have missing values.")
        return mis_val_tab_rename_cols
```

```
In [11]: missing_values = mis_values(df1)
```

Your selected dataframe has 38 columns.

There are 21 cols that have missing values.

```
In [12]: print(missing_values)
```

	Missing Values	% of Total Missing Values
gross	884	17.5
budget	492	9.8
aspect_ratio	329	6.5
content_rating	303	6.0
plot_keywords	153	3.0
title_year	108	2.1
director_name	104	2.1
director_facebook_likes	104	2.1
num_critic_for_reviews	50	1.0
actor_3_name	23	0.5
actor_3_facebook_likes	23	0.5
num_user_for_reviews	21	0.4
color	19	0.4
duration	15	0.3
facenumber_in_poster	13	0.3
actor_2_name	13	0.3
actor_2_facebook_likes	13	0.3
language	12	0.2
actor_1_name	7	0.1

actor_1_facebook_likes	7	0.1
country	5	0.1

Instead of dropping the rows, I used the median imputation because it's maintain the distribution of the variable.

```
In [13]: # fill missing values with median column values
df1 = df1.fillna(df1.median())
```

```
In [14]: df1
```

```
Out[14]:
```

	id	stock_market_idx	days_since_last_tweet	pre_screen_viewers	\
0	1	1102	67	18	
1	2	1117	70	18	
2	3	1000	90	11	
3	4	1007	35	10	
4	5	1128	85	20	
5	6	1037	75	20	
6	7	1021	80	12	
7	8	1133	30	16	
8	9	1186	14	11	
9	10	1016	3	10	
10	11	1142	99	13	
11	12	1197	35	17	
12	13	1105	17	12	
13	14	1145	8	10	
14	15	1010	14	15	
15	16	1146	89	14	
16	17	1188	12	18	
17	18	1156	17	12	
18	19	1185	59	12	
19	20	1127	77	11	
20	21	1051	36	19	
21	22	1015	61	19	
22	23	1200	75	14	
23	24	1163	45	10	
24	25	1042	73	19	
25	26	1142	96	16	
26	27	1086	28	19	
27	28	1114	41	13	
28	29	1197	66	16	
29	30	1196	23	19	
...	...	...	...	...	...
5013	5014	1081	57	10	
5014	5015	1174	77	12	
5015	5016	1186	15	19	
5016	5017	1136	59	11	
5017	5018	1110	19	10	

5018	5019	1040	66	19
5019	5020	1183	43	19
5020	5021	1155	61	11
5021	5022	1027	96	11
5022	5023	1190	13	13
5023	5024	1138	77	12
5024	5025	1055	13	16
5025	5026	1125	97	12
5026	5027	1170	58	10
5027	5028	1157	98	16
5028	5029	1122	69	19
5029	5030	1056	49	20
5030	5031	1193	97	18
5031	5032	1196	31	15
5032	5033	1026	76	20
5033	5034	1177	84	11
5034	5035	1060	59	16
5035	5036	1167	2	10
5036	5037	1062	45	12
5037	5038	1013	57	13
5038	5039	1131	42	19
5039	5040	1016	26	11
5040	5041	1146	33	18
5041	5042	1015	31	11
5042	5043	1048	9	19

	characters_per_longest_review	priority	\
0	1181	4	
1	1196	4	
2	1125	4	
3	1127	4	
4	1072	4	
5	1121	4	
6	1129	4	
7	1164	4	
8	1076	4	
9	1040	4	
10	1191	4	
11	1154	4	
12	1019	4	
13	1045	4	
14	1135	4	
15	1163	4	
16	1066	4	
17	1020	4	
18	1165	4	
19	1142	4	
20	1057	4	

21	1147	4
22	1103	4
23	1092	4
24	1005	4
25	1134	4
26	1035	4
27	1125	4
28	1158	4
29	1022	4
...	...	...
5013	1171	4
5014	1091	4
5015	1101	4
5016	1172	4
5017	1012	4
5018	1170	4
5019	1074	4
5020	1140	4
5021	1002	4
5022	1105	4
5023	1064	4
5024	1097	4
5025	1086	4
5026	1003	4
5027	1095	4
5028	1108	4
5029	1146	4
5030	1027	4
5031	1036	4
5032	1177	4
5033	1159	4
5034	1133	4
5035	1152	4
5036	1059	4
5037	1187	4
5038	1073	4
5039	1002	4
5040	1021	4
5041	1020	4
5042	1081	4

	longest_facebook_comment_review_char	color \
0	250	Color
1	740	Color
2	1779	Color
3	1074	Color
4	813	NaN
5	508	Color



6	1189	Color
7	842	Color
8	1860	Color
9	832	Color
10	637	Color
11	1083	Color
12	429	Color
13	1506	Color
14	1389	Color
15	87	Color
16	829	Color
17	1256	Color
18	729	Color
19	737	Color
20	557	Color
21	659	Color
22	1742	Color
23	1770	Color
24	1028	Color
25	590	Color
26	390	Color
27	760	Color
28	1413	Color
29	91	Color
...	...	...
5013	1530	Color
5014	1897	Color
5015	999	Black and White
5016	1724	Color
5017	1022	Color
5018	1434	Color
5019	675	Color
5020	625	NaN
5021	1479	Color
5022	221	Black and White
5023	923	Color
5024	1828	Color
5025	1342	Color
5026	971	Color
5027	1145	Color
5028	311	Black and White
5029	1489	Color
5030	684	Color
5031	1121	Color
5032	1091	Color
5033	1365	Color
5034	181	Color
5035	1497	Color

5036	373	Color
5037	766	Color
5038	10	Color
5039	539	Color
5040	1558	Color
5041	1152	Color
5042	1587	Color

	director_name	num_critic_for_reviews	...	country \
0	James Cameron	723.0	...	USA
1	Gore Verbinski	302.0	...	USA
2	Sam Mendes	602.0	...	UK
3	Christopher Nolan	813.0	...	USA
4	Doug Walker	110.0	...	NaN
5	Andrew Stanton	462.0	...	USA
6	Sam Raimi	392.0	...	USA
7	Nathan Greno	324.0	...	USA
8	Joss Whedon	635.0	...	USA
9	David Yates	375.0	...	UK
10	Zack Snyder	673.0	...	USA
11	Bryan Singer	434.0	...	USA
12	Marc Forster	403.0	...	UK
13	Gore Verbinski	313.0	...	USA
14	Gore Verbinski	450.0	...	USA
15	Zack Snyder	733.0	...	USA
16	Andrew Adamson	258.0	...	USA
17	Joss Whedon	703.0	...	USA
18	Rob Marshall	448.0	...	USA
19	Barry Sonnenfeld	451.0	...	USA
20	Peter Jackson	422.0	...	New Zealand
21	Marc Webb	599.0	...	USA
22	Ridley Scott	343.0	...	USA
23	Peter Jackson	509.0	...	USA
24	Chris Weitz	251.0	...	USA
25	Peter Jackson	446.0	...	New Zealand
26	James Cameron	315.0	...	USA
27	Anthony Russo	516.0	...	USA
28	Peter Berg	377.0	...	USA
29	Colin Trevorrow	644.0	...	USA
...	...	...	...	...
5013	Eric Eason	28.0	...	USA
5014	Uwe Boll	58.0	...	Canada
5015	Richard Linklater	61.0	...	USA
5016	Joseph Mazzella	110.0	...	USA
5017	Travis Legge	1.0	...	USA
5018	Alex Kendrick	5.0	...	USA
5019	Marcus Nispel	43.0	...	USA
5020	Brandon Landers	110.0	...	USA

5021	Jay Duplass	51.0	...	USA
5022	Jim Chuchu	6.0	...	Kenya
5023	Daryl Wein	22.0	...	USA
5024	Jason Trost	42.0	...	USA
5025	John Waters	73.0	...	USA
5026	Olivier Assayas	81.0	...	France
5027	Jafar Panahi	64.0	...	Iran
5028	Ivan Kavanagh	12.0	...	Ireland
5029	Kiyoshi Kurosawa	78.0	...	Japan
5030	Tadeo Garcia	110.0	...	USA
5031	Thomas L. Phillips	13.0	...	USA
5032	Ash Baron-Cohen	10.0	...	USA
5033	Shane Carruth	143.0	...	USA
5034	Neill Dela Llana	35.0	...	Philippines
5035	Robert Rodriguez	56.0	...	USA
5036	Anthony Vallone	110.0	...	USA
5037	Edward Burns	14.0	...	USA
5038	Scott Smith	1.0	...	Canada
5039	NaN	43.0	...	USA
5040	Benjamin Roberds	13.0	...	USA
5041	Daniel Hsia	14.0	...	USA
5042	Jon Gunn	43.0	...	USA

	content_rating	website_score	budget	weighted_budget	title_year \
0	PG-13	7.9	237000000.0	236999000	2009.0
1	PG-13	7.1	300000000.0	299999000	2007.0
2	PG-13	6.8	245000000.0	244999000	2015.0
3	PG-13	8.5	250000000.0	249999000	2012.0
4	NaN	7.1	20000000.0	-1000	2005.0
5	PG-13	6.6	263700000.0	263699000	2012.0
6	PG-13	6.2	258000000.0	257999000	2007.0
7	PG	7.8	260000000.0	259999000	2010.0
8	PG-13	7.5	250000000.0	249999000	2015.0
9	PG	7.5	250000000.0	249999000	2009.0
10	PG-13	6.9	250000000.0	249999000	2016.0
11	PG-13	6.1	209000000.0	208999000	2006.0
12	PG-13	6.7	200000000.0	199999000	2008.0
13	PG-13	7.3	225000000.0	224999000	2006.0
14	PG-13	6.5	215000000.0	214999000	2013.0
15	PG-13	7.2	225000000.0	224999000	2013.0
16	PG	6.6	225000000.0	224999000	2008.0
17	PG-13	8.1	220000000.0	219999000	2012.0
18	PG-13	6.7	250000000.0	249999000	2011.0
19	PG-13	6.8	225000000.0	224999000	2012.0
20	PG-13	7.5	250000000.0	249999000	2014.0
21	PG-13	7.0	230000000.0	229999000	2012.0
22	PG-13	6.7	200000000.0	199999000	2010.0
23	PG-13	7.9	225000000.0	224999000	2013.0

24	PG-13	6.1	180000000.0	179999000	2007.0
25	PG-13	7.2	207000000.0	206999000	2005.0
26	PG-13	7.7	200000000.0	199999000	1997.0
27	PG-13	8.2	250000000.0	249999000	2016.0
28	PG-13	5.9	209000000.0	208999000	2012.0
29	PG-13	7.0	150000000.0	149999000	2015.0
...	...	...	...	...	...
5013	NaN	7.0	24000.0	23000	2002.0
5014	R	6.3	20000000.0	-1000	2009.0
5015	R	7.1	23000.0	22000	1991.0
5016	NaN	4.8	25000.0	24000	2015.0
5017	NaN	3.3	22000.0	21000	2013.0
5018	NaN	6.9	20000.0	19000	2003.0
5019	R	4.6	20000000.0	-1000	2015.0
5020	NaN	3.0	17350.0	16350	2011.0
5021	R	6.6	15000.0	14000	2005.0
5022	NaN	7.4	15000.0	14000	2014.0
5023	NaN	6.2	15000.0	14000	2009.0
5024	Unrated	4.0	20000.0	19000	2011.0
5025	NC-17	6.1	10000.0	9000	1972.0
5026	R	6.9	4500.0	3500	2004.0
5027	Not Rated	7.5	10000.0	9000	2000.0
5028	NaN	6.7	10000.0	9000	2007.0
5029	NaN	7.4	1000000.0	999000	1997.0
5030	NaN	6.1	20000000.0	-1000	2004.0
5031	NaN	5.4	200000.0	199000	2012.0
5032	NaN	6.4	20000000.0	-1000	1995.0
5033	PG-13	7.0	7000.0	6000	2004.0
5034	Not Rated	6.3	7000.0	6000	2005.0
5035	R	6.9	7000.0	6000	1992.0
5036	PG-13	7.8	3250.0	2250	2005.0
5037	Not Rated	6.4	9000.0	8000	2011.0
5038	NaN	7.7	20000000.0	-1000	2013.0
5039	TV-14	7.5	20000000.0	-1000	2005.0
5040	NaN	6.3	1400.0	400	2013.0
5041	PG-13	6.3	20000000.0	-1000	2012.0
5042	PG	6.6	1100.0	100	2004.0

	actor_2_facebook_likes	aspect_ratio	movie_facebook_likes	imdb_score
0	936.0	1.78	33000	7.9
1	5000.0	2.35	0	7.1
2	393.0	2.35	85000	6.8
3	23000.0	2.35	164000	8.5
4	12.0	2.35	0	7.1
5	632.0	2.35	24000	6.6
6	11000.0	2.35	0	6.2
7	553.0	1.85	29000	7.8
8	21000.0	2.35	118000	7.5

9	11000.0	2.35	10000	7.5
10	4000.0	2.35	197000	6.9
11	10000.0	2.35	0	6.1
12	412.0	2.35	0	6.7
13	5000.0	2.35	5000	7.3
14	2000.0	2.35	48000	6.5
15	3000.0	2.35	118000	7.2
16	216.0	2.35	0	6.6
17	21000.0	1.85	123000	8.1
18	11000.0	2.35	58000	6.7
19	816.0	1.85	40000	6.8
20	972.0	2.35	65000	7.5
21	10000.0	2.35	56000	7.0
22	882.0	2.35	17000	6.7
23	972.0	2.35	83000	7.9
24	6000.0	2.35	0	6.1
25	919.0	2.35	0	7.2
26	14000.0	2.35	26000	7.7
27	19000.0	2.35	72000	8.2
28	10000.0	2.35	44000	5.9
29	2000.0	2.00	150000	7.0
...	...	...	...	...
5013	46.0	1.78	61	7.0
5014	918.0	2.35	0	6.3
5015	0.0	1.37	2000	7.1
5016	25.0	2.35	33	4.8
5017	184.0	1.78	200	3.3
5018	49.0	1.85	725	6.9
5019	512.0	1.85	0	4.6
5020	19.0	2.35	33	3.0
5021	224.0	2.35	297	6.6
5022	19.0	2.35	45	7.4
5023	212.0	2.35	324	6.2
5024	91.0	2.35	835	4.0
5025	143.0	1.37	0	6.1
5026	133.0	2.35	171	6.9
5027	0.0	1.85	697	7.5
5028	5.0	1.33	105	6.7
5029	13.0	1.85	817	7.4
5030	20.0	2.35	22	6.1
5031	98.0	16.00	424	5.4
5032	194.0	2.35	20	6.4
5033	45.0	1.85	19000	7.0
5034	0.0	2.35	74	6.3
5035	20.0	1.37	0	6.9
5036	44.0	2.35	4	7.8
5037	205.0	2.35	413	6.4
5038	470.0	2.35	84	7.7

5039	593.0	16.00	32000	7.5
5040	0.0	2.35	16	6.3
5041	719.0	2.35	660	6.3
5042	23.0	1.85	456	6.6

[5043 rows x 38 columns]

In [15]: df1.dtypes.value\_counts()

Out[15]: float64 15  
object 12  
int64 11  
dtype: int64

We do! Let's use the "describe" method to find them, amongst other interesting information

In [16]: df1.describe()

Out[16]:

	id	stock_market_idx	days_since_last_tweet	\
count	5043.000000	5043.000000	5043.000000	
mean	2522.000000	1101.16042	49.908190	
std	1455.933034	58.48476	28.432368	
min	1.000000	1000.00000	1.000000	
25%	1261.500000	1051.00000	26.000000	
50%	2522.000000	1101.00000	49.000000	
75%	3782.500000	1153.00000	74.000000	
max	5043.000000	1200.00000	99.000000	

	pre_screen_viewers	characters_per_longest_review	priority	\
count	5043.000000	5043.000000	5043.0	
mean	14.979576	1100.040849	4.0	
std	3.163246	57.299452	0.0	
min	10.000000	1000.000000	4.0	
25%	12.000000	1051.000000	4.0	
50%	15.000000	1099.000000	4.0	
75%	18.000000	1149.000000	4.0	
max	20.000000	1200.000000	4.0	

	longest_facebook_comment_review_char	num_critic_for_reviews	\
count	5043.000000	5043.000000	
mean	962.646242	139.894904	
std	541.710282	121.034214	
min	6.000000	1.000000	
25%	511.000000	50.000000	
50%	962.000000	110.000000	
75%	1424.000000	194.000000	
max	1900.000000	813.000000	

	duration	director_facebook_likes	...	\
--	----------	-------------------------	-----	---

count	5043.000000	5043.000000	...
mean	107.188578	673.362086	...
std	25.160972	2785.636586	...
min	7.000000	0.000000	...
25%	93.000000	7.000000	...
50%	103.000000	49.000000	...
75%	118.000000	189.000000	...
max	511.000000	23000.000000	...

	facenumber_in_poster	num_user_for_reviews	website_score \
count	5043.000000	5043.000000	5043.000000
mean	1.370216	272.284553	6.559925
std	2.011066	377.269873	8.433695
min	0.000000	1.000000	1.600000
25%	0.000000	65.000000	5.800000
50%	1.000000	156.000000	6.600000
75%	2.000000	324.000000	7.200000
max	43.000000	5060.000000	600.000000

	budget	weighted_budget	title_year	actor_2_facebook_likes \
count	5.043000e+03	5.043000e+03	5043.000000	5043.000000
mean	3.782554e+07	3.587332e+07	2002.531033	1649.030339
std	1.958882e+08	1.961555e+08	12.359307	4037.579765
min	2.180000e+02	-1.000000e+03	1916.000000	0.000000
25%	7.000000e+06	2.999000e+06	1999.000000	281.000000
50%	2.000000e+07	1.499900e+07	2005.000000	595.000000
75%	4.000000e+07	3.999900e+07	2011.000000	918.000000
max	1.221550e+10	1.221550e+10	2045.000000	137000.000000

	aspect_ratio	movie_facebook_likes	imdb_score
count	5043.000000	5043.000000	5043.000000
mean	2.228858	7525.964505	6.559925
std	1.339542	19320.445110	8.433695
min	1.180000	0.000000	1.600000
25%	1.850000	0.000000	5.800000
50%	2.350000	166.000000	6.600000
75%	2.350000	3000.000000	7.200000
max	16.000000	349000.000000	600.000000

[8 rows x 26 columns]

I dropped the columns because these columns had high cardinality or many levels and also many zero values.

```
In [17]: # Drop extraneous columns
```

```
col = ['id', 'priority', 'director_name', 'actor_2_name', 'color', 'actor_1_name', 'a
'director_name', 'facenumber_in_poster', 'movie_facebook_likes', 'website_score
'movie_imdb_link', 'content_rating', 'language', 'plot_keywords', 'cast_total_
```

```

        'director_facebook_likes', 'actor_2_facebook_likes', 'actor_3_facebook_likes',
        'aspect_ratio', 'website_score', 'genres', 'actor_1_facebook_likes']
df1.drop(col, axis=1, inplace=True)

```

```
In [18]: df1.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5043 entries, 0 to 5042
Data columns (total 13 columns):
stock_market_idx          5043 non-null int64
days_since_last_tweet    5043 non-null int64
pre_screen_viewers        5043 non-null int64
characters_per_longest_review  5043 non-null int64
longest_facebook_comment_review_char  5043 non-null int64
num_critic_for_reviews    5043 non-null float64
duration                  5043 non-null float64
gross                     5043 non-null float64
num_voted_users           5043 non-null int64
made_up_column            5043 non-null float64
num_user_for_reviews      5043 non-null float64
budget                    5043 non-null float64
imdb_score                5043 non-null float64
dtypes: float64(7), int64(6)
memory usage: 551.6 KB

```

## <h2>4. Data Transformation</h2>

To normalize the above variables, I used log transformation.

```

In [19]: df1['budget'] = np.log(df1.budget)
        df1['imdb_score'] = np.log(df1.imdb_score)
        df1['made_up_column'] = np.log(df1.imdb_score)
        df1['stock_market_idx'] = np.log(df1.stock_market_idx)
        df1['days_since_last_tweet'] = np.log(df1.days_since_last_tweet)
        df1['pre_screen_viewers'] = np.log(df1.pre_screen_viewers)
        df1['characters_per_longest_review'] = np.log(df1.characters_per_longest_review)
        df1['longest_facebook_comment_review_char'] = np.log(df1.longest_facebook_comment_rev
        df1['num_critic_for_reviews'] = np.log(df1.num_critic_for_reviews)
        df1['duration'] = np.log(df1.duration)
        df1['gross'] = np.log(df1.gross)
        df1['num_voted_users'] = np.log(df1.num_voted_users)
        df1['num_user_for_reviews'] = np.log(df1.num_user_for_reviews)

```

```
In [20]: df1.describe()
```

```

Out[20]:
      stock_market_idx  days_since_last_tweet  pre_screen_viewers  \
count          5043.000000          5043.000000          5043.000000
mean              7.002705              3.630552              2.683450

```



std	0.053234	0.913946	0.218117
min	6.907755	0.000000	2.302585
25%	6.957497	3.258097	2.484907
50%	7.003974	3.891820	2.708050
75%	7.050123	4.304065	2.890372
max	7.090077	4.595120	2.995732

	characters_per_longest_review	longest_facebook_comment_review_char	\
count	5043.000000	5043.000000	
mean	7.001743	6.583890	
std	0.052173	0.944422	
min	6.907755	1.791759	
25%	6.957497	6.236370	
50%	7.002156	6.869014	
75%	7.046647	7.261225	
max	7.090077	7.549609	

	num_critic_for_reviews	duration	gross	num_voted_users	\
count	5043.000000	5043.000000	5043.000000	5043.000000	
mean	4.467420	4.647760	16.501699	10.096277	
std	1.163810	0.242031	2.125449	1.990129	
min	0.000000	1.945910	5.087596	1.609438	
25%	3.912023	4.532599	15.950977	9.058761	
50%	4.700480	4.634729	17.054875	10.444619	
75%	5.267858	4.770685	17.754313	11.475317	
max	6.700731	6.236370	20.449494	14.340099	

	made_up_column	num_user_for_reviews	budget	imdb_score
count	5043.000000	5043.000000	5043.000000	5043.000000
mean	0.605564	4.871133	16.478166	1.845742
std	0.126912	1.392721	1.635744	0.210021
min	-0.755015	0.000000	5.384495	0.470004
25%	0.564096	4.174387	15.761421	1.757858
50%	0.635025	5.049856	16.811243	1.887070
75%	0.680103	5.780744	17.504390	1.974081
max	1.855818	8.529122	23.225971	6.396930

In [21]: df1.profile\_report()

<IPython.lib.display.IFrame at 0x1fa2e8a0080>

Out[21]:

Some of the columns have missing values. We can deal with this in a few different ways. The simplest solution is to remove them, though we lose many examples in doing so. Alternatively, we could impute the values, replacing the NaN values with an average (mean or median).

For the purpose of this simple notebook, the variable num\_user\_for\_reviews has 51 zeros which was replaced by median and also remaining numerical variable are imputed by median imputation.

```

In [22]: nonzero_median = df1[ df1.num_user_for_reviews != 0 ].median()

In [23]: df1.loc[ df1.num_user_for_reviews == 0, "num_user_for_reviews" ] = nonzero_median

In [24]: # fill missing values with median column values
         df1 = df1.fillna(df1.median())

In [25]: df1.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5043 entries, 0 to 5042
Data columns (total 13 columns):
stock_market_idx          5043 non-null float64
days_since_last_tweet    5043 non-null float64
pre_screen_viewers        5043 non-null float64
characters_per_longest_review  5043 non-null float64
longest_facebook_comment_review_char  5043 non-null float64
num_critic_for_reviews    5043 non-null float64
duration                  5043 non-null float64
gross                     5043 non-null float64
num_voted_users           5043 non-null float64
made_up_column            5043 non-null float64
num_user_for_reviews      5043 non-null float64
budget                    5043 non-null float64
imdb_score                5043 non-null float64
dtypes: float64(13)
memory usage: 551.6 KB

```

## <h2>5. Data Modeling</h2>

### 5.1. Data Splitting

#### 0.0.2 Test & Train Split

The purpose of splitting the data is to be able to assess the quality of a predictive model when it is used on unseen data. When training, you will try to build a model that fits to the data as closely as possible, to be able to most accurately make a prediction. However, without a test set you run the risk of overfitting - the model works very well for the data it has seen but not for new data.

The split ratio is often debated and in practice you might split your data into three sets: train, validation and test. You would use the training data to understand which classifier you wish to use; the validation set to test on whilst tweaking parameters; and the test set to get an understanding of how your final model would work in practice. Furthermore, there are techniques such as K-Fold cross validation that also help to reduce bias.

For the purpose of this demonstration, we will only be randomly splitting our data into test and train, with a 80/20 split.

We import the required library from scikit-learn, `train_test_split`

```

In [26]: X = df1.iloc[:,0:12].values
         y = df1.iloc[:,12:13].values

```

```
In [27]: X_trainset, X_testset, y_trainset, y_testset = train_test_split(X, y, test_size=0.20,
```

```
In [28]: print(X_trainset.shape)
         print(y_trainset.shape)
```

```
(4034, 12)
```

```
(4034, 1)
```

```
In [29]: print(X_testset.shape)
         print(y_testset.shape)
```

```
(1009, 12)
```

```
(1009, 1)
```

## <h2> Parametric Machine Learning Algorithm</h2>

### 5.2. Using Linear Regression

Linear regression attempts to fit a straight hyperplane to your dataset that is closest to all data points. It is most suitable when there are linear relationships between the variables in the dataset.

```
In [30]: reg = LinearRegression()
         reg.fit(X, y)
```

```
Out[30]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [31]: y_pred1 = reg.predict(X_testset)
```

#### 5.2.1. Performance Metrics of Linear Regression

We wish to understand how good our model is; there are a few different metrics we can use. We will evaluate mean squared error (MSE) and mean absolute error (MAE)

We import scikit-learn's mean squared error and scikit-learn's mean absolute error

```
In [32]: from sklearn import metrics
         print('Mean Absolute Error:', metrics.mean_absolute_error(y_testset, y_pred1))
         print('Mean Squared Error:', metrics.mean_squared_error(y_testset, y_pred1))
         print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_testset, y_pred1)))
```

```
Mean Absolute Error: 0.01813151995508253
```

```
Mean Squared Error: 0.0014731432685910185
```

```
Root Mean Squared Error: 0.03838154854342147
```

### 5.3. Using Ridge Regression

```
In [33]: from sklearn.model_selection import GridSearchCV
         from sklearn.linear_model import Ridge
```

```

In [34]: ridge = Ridge()
         parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}
         ridge_regressor = GridSearchCV(ridge, parameters, scoring='mean_squared_error')

         ridge_regressor.fit(X, y)

Out[34]: GridSearchCV(cv=None, error_score='raise',
                    estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
                    normalize=False, random_state=None, solver='auto', tol=0.001),
                    fit_params=None, iid=True, n_jobs=1,
                    param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1, 5, 10, 20]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                    scoring='mean_squared_error', verbose=0)

In [35]: print(ridge_regressor.best_params_)
         print(ridge_regressor.best_score_)

{'alpha': 1}
-0.00248930108763342

```

In this case, the optimal value for alpha is 1, and the negative MSE is -0.0024893.

#### 5.4. Using LASSO Regression

```

In [36]: from sklearn.linear_model import Lasso

In [37]: lasso = Lasso()
         parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}
         lasso_regressor = GridSearchCV(lasso, parameters, scoring='mean_squared_error')

         lasso_regressor.fit(X, y)

Out[37]: GridSearchCV(cv=None, error_score='raise',
                    estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=1000,
                    normalize=False, positive=False, precompute=False, random_state=None,
                    selection='cyclic', tol=0.0001, warm_start=False),
                    fit_params=None, iid=True, n_jobs=1,
                    param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1, 5, 10, 20]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                    scoring='mean_squared_error', verbose=0)

In [38]: print(lasso_regressor.best_params_)
         print(lasso_regressor.best_score_)

{'alpha': 0.0001}
-0.0024895672740462694

```

In this case, the optimal value for alpha is 0.0001, and the negative MSE is -0.0024895.

Note: After use linear, lasso, and ridge regression. We have seen that ridge is the best fitting method, with a regularization value of 1.

## <h2>Nonparametric ML Algorithms</h2>

### 5.5. Using KNeighborsRegressor

```
In [39]: clf=KNeighborsRegressor(5)
         clf.fit(X_trainset,y_trainset)
         y_pred=clf.predict(X_testset)
         print(mean_squared_error(y_testset,y_pred))
```

0.031100714262353453


```
In [42]: from sklearn import neighbors
         rmse_val = [] #to store rmse values for different k
         for K in range(20):
             K = K+1
             model = neighbors.KNeighborsRegressor(n_neighbors = K)

             model.fit(X_trainset, y_trainset) #fit the model
             pred=model.predict(X_testset) #make prediction on test set
             error = sqrt(mean_squared_error(y_testset,pred)) #calculate rmse
             rmse_val.append(error) #store rmse values
             print('RMSE value for k= ' , K , 'is:', error)
```

```
RMSE value for k= 1 is: 0.21315240080835277
RMSE value for k= 2 is: 0.1908013268343305
RMSE value for k= 3 is: 0.18098856364713534
RMSE value for k= 4 is: 0.17803795746340828
RMSE value for k= 5 is: 0.17635394597896994
RMSE value for k= 6 is: 0.17416175135818351
RMSE value for k= 7 is: 0.1731860406522612
RMSE value for k= 8 is: 0.17230452241536612
RMSE value for k= 9 is: 0.17302494104578212
RMSE value for k= 10 is: 0.17256339138196394
RMSE value for k= 11 is: 0.17266004948251795
RMSE value for k= 12 is: 0.1728858069932387
RMSE value for k= 13 is: 0.17274587714990983
RMSE value for k= 14 is: 0.1730896798547398
RMSE value for k= 15 is: 0.17308010745515268
RMSE value for k= 16 is: 0.1730563602443598
RMSE value for k= 17 is: 0.17335005398973438
RMSE value for k= 18 is: 0.17323566744883404
RMSE value for k= 19 is: 0.1736963168112529
RMSE value for k= 20 is: 0.17340505276908819
```

```
In [43]: #plotting the rmse values against k values
         curve = pd.DataFrame(rmse_val) #elbow curve
         curve.plot()
```

Out [43]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fa2e148588>



Comparing with other Number of clusters,  $k=8$  and RMSE is 0.172 is better.

#### 5.6. Using of Random Forest

```
In [44]: regressor = RandomForestRegressor(n_estimators=20, random_state=0)
         regressor.fit(X_trainset, y_trainset)
         y_pred = regressor.predict(X_testset)
```

```
In [45]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_testset, y_pred))
         print('Mean Squared Error:', metrics.mean_squared_error(y_testset, y_pred))
         print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_testset, y_pred)))
```

Mean Absolute Error: 0.00044979607149377013

Mean Squared Error: 4.4625280188515924e-05

Root Mean Squared Error: 0.006680215579494117

When estimator is 20 then RMSE is 0.007 and MSE is 4.46.

Note:

After comparing non-parametric models RMSE, Random forest performed better.

### 0.0.3 Stepwise Selection

Used backward selection method for multiple regression to find better R-square and all p-value of variable should be significantly significant (less than 0.05)

```
In [46]: X = np.column_stack((df1['budget'], df1['stock_market_idx'], df1['days_since_last_tweet'],
                             df1['characters_per_longest_review'], df1['longest_facebook_comment'],
                             df1['num_critic_for_reviews'], df1['duration'], df1['gross'], df1['gross_per_m'],
                             df1['made_up_column'], df1['num_user_for_reviews']))
y = df1['imdb_score']
X2 = sm.add_constant(X)
est = sm.OLS(y, X2)
est2 = est.fit()
print(est2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  imdb_score    R-squared:                0.948
Model:                            OLS        Adj. R-squared:            0.948
Method:                    Least Squares    F-statistic:                7691.
Date:                Sat, 20 Jul 2019    Prob (F-statistic):          0.00
Time:                17:39:57    Log-Likelihood:            8184.7
No. Observations:                5043    AIC:                      -1.634e+04
Df Residuals:                    5030    BIC:                      -1.626e+04
Df Model:                        12
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1.0315	0.129	7.993	0.000	0.778	1.284
x1	-0.0032	0.001	-6.456	0.000	-0.004	-0.002
x2	-0.0086	0.013	-0.678	0.498	-0.033	0.016
x3	-0.0001	0.001	-0.139	0.890	-0.002	0.001
x4	-0.0003	0.003	-0.106	0.915	-0.006	0.006
x5	-0.0137	0.013	-1.059	0.290	-0.039	0.012
x6	-0.0008	0.001	-1.065	0.287	-0.002	0.001
x7	-0.0059	0.001	-5.906	0.000	-0.008	-0.004
x8	0.0068	0.003	2.282	0.023	0.001	0.013
x9	-5.332e-05	0.000	-0.143	0.886	-0.001	0.001
x10	0.0041	0.001	5.531	0.000	0.003	0.006
x11	1.5940	0.006	279.342	0.000	1.583	1.605
x12	0.0037	0.001	3.635	0.000	0.002	0.006

```

=====
Omnibus:                13727.176    Durbin-Watson:                1.984
Prob(Omnibus):                0.000    Jarque-Bera (JB):            583187610.018
Skew:                33.242    Prob(JB):                0.00
Kurtosis:                1667.636    Cond. No.                5.72e+03
=====

```

#### Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.72e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [47]: col = ['pre_screen_viewers']
         df1.drop(col, axis=1, inplace=True)

In [48]: X = np.column_stack((df1['budget'], df1['stock_market_idx'], df1['days_since_last_tweet'],
                             df1['characters_per_longest_review'], df1['longest_facebook_comment'],
                             df1['num_critic_for_reviews'], df1['duration'], df1['gross'], df1['gross_per_sq_meter'],
                             df1['made_up_column'], df1['num_user_for_reviews']))

         y = df1['imdb_score']
         X2 = sm.add_constant(X)
         est = sm.OLS(y, X2)
         est2 = est.fit()
         print(est2.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                  imdb_score    R-squared:                0.948
Model:                            OLS        Adj. R-squared:            0.948
Method:                    Least Squares    F-statistic:                8391.
Date:                Sat, 20 Jul 2019    Prob (F-statistic):          0.00
Time:                17:39:59    Log-Likelihood:            8184.7
No. Observations:                5043    AIC:                      -1.635e+04
Df Residuals:                    5031    BIC:                      -1.627e+04
Df Model:                        11
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1.0307	0.129	8.001	0.000	0.778	1.283
x1	-0.0032	0.001	-6.456	0.000	-0.004	-0.002
x2	-0.0086	0.013	-0.678	0.498	-0.033	0.016
x3	-0.0001	0.001	-0.140	0.889	-0.002	0.001
x4	-0.0137	0.013	-1.060	0.289	-0.039	0.012
x5	-0.0008	0.001	-1.066	0.287	-0.002	0.001
x6	-0.0059	0.001	-5.906	0.000	-0.008	-0.004
x7	0.0068	0.003	2.282	0.023	0.001	0.013
x8	-5.33e-05	0.000	-0.143	0.886	-0.001	0.001
x9	0.0041	0.001	5.533	0.000	0.003	0.006
x10	1.5940	0.006	279.385	0.000	1.583	1.605
x11	0.0037	0.001	3.634	0.000	0.002	0.006

```

=====
Omnibus:                13727.412    Durbin-Watson:                1.984
Prob(Omnibus):                0.000    Jarque-Bera (JB):            583262636.992
Skew:                33.244    Prob(JB):                0.00
Kurtosis:                1667.743    Cond. No.                5.69e+03
=====

```



Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.69e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [49]: col = ['gross']
         df1.drop(col, axis=1, inplace=True)

In [50]: X = np.column_stack((df1['budget'], df1['stock_market_idx'], df1['days_since_last_tweet'],
                             df1['characters_per_longest_review'], df1['longest_facebook_comment'],
                             df1['num_critic_for_reviews'], df1['duration'], df1['num_voted_users'],
                             df1['made_up_column'], df1['num_user_for_reviews']))
         y = df1['imdb_score']
         X2 = sm.add_constant(X)
         est = sm.OLS(y, X2)
         est2 = est.fit()
         print(est2.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          imdb_score    R-squared:                0.948
Model:                  OLS          Adj. R-squared:            0.948
Method:                 Least Squares   F-statistic:             9232.
Date:                  Sat, 20 Jul 2019   Prob (F-statistic):       0.00
Time:                  17:40:00          Log-Likelihood:           8184.7
No. Observations:      5043             AIC:                   -1.635e+04
Df Residuals:          5032             BIC:                   -1.628e+04
Df Model:               10
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	1.0306	0.129	8.002	0.000	0.778	1.283
x1	-0.0033	0.000	-6.673	0.000	-0.004	-0.002
x2	-0.0086	0.013	-0.681	0.496	-0.033	0.016
x3	-0.0001	0.001	-0.138	0.890	-0.002	0.001
x4	-0.0137	0.013	-1.063	0.288	-0.039	0.012
x5	-0.0008	0.001	-1.067	0.286	-0.002	0.001
x6	-0.0058	0.001	-5.953	0.000	-0.008	-0.004
x7	0.0069	0.003	2.291	0.022	0.001	0.013
x8	0.0041	0.001	5.572	0.000	0.003	0.006
x9	1.5940	0.006	280.066	0.000	1.583	1.605
x10	0.0037	0.001	3.667	0.000	0.002	0.006

```
=====
Omnibus:                13726.670    Durbin-Watson:           1.984
Prob(Omnibus):           0.000      Jarque-Bera (JB):        582997176.363
Skew:                    33.239      Prob(JB):                0.00
=====
```

Kurtosis: 1667.364 Cond. No. 4.71e+03

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 4.71e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [51]: col = ['days_since_last_tweet']
df1.drop(col, axis=1, inplace=True)
```

```
In [52]: X = np.column_stack((df1['budget'], df1['stock_market_idx'],
                             df1['characters_per_longest_review'], df1['longest_facebook_comm
                             df1['num_critic_for_reviews'], df1['duration'], df1['num_voted_us
                             df1['made_up_column'], df1['num_user_for_reviews']]))
y = df1['imdb_score']
X2 = sm.add_constant(X)
est = sm.OLS(y, X2)
est2 = est.fit()
print(est2.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          imdb_score    R-squared:                0.949
Model:                  OLS          Adj. R-squared:            0.949
Method:                 Least Squares   F-statistic:            1.043e+04
Date:                  Sat, 20 Jul 2019   Prob (F-statistic):       0.00
Time:                  17:40:01          Log-Likelihood:           8223.5
No. Observations:      5043             AIC:                   -1.643e+04
Df Residuals:          5033             BIC:                   -1.636e+04
Df Model:               9
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	1.0974	0.128	8.583	0.000	0.847	1.348
x1	-0.0028	0.000	-6.085	0.000	-0.004	-0.002
x2	-0.0113	0.013	-0.897	0.370	-0.036	0.013
x3	-0.0151	0.013	-1.177	0.239	-0.040	0.010
x4	-0.0007	0.001	-1.023	0.307	-0.002	0.001
x5	-0.0035	0.001	-4.009	0.000	-0.005	-0.002
x6	0.0036	0.003	1.209	0.227	-0.002	0.009
x7	6.57e-08	6.27e-09	10.472	0.000	5.34e-08	7.8e-08
x8	1.5878	0.006	279.653	0.000	1.577	1.599
x9	0.0033	0.001	3.751	0.000	0.002	0.005

```
=====
Omnibus:                13897.771    Durbin-Watson:           1.989
```

```

Prob(Omnibus):                0.000    Jarque-Bera (JB):                640334751.739
Skew:                        34.339    Prob(JB):                      0.00
Kurtosis:                    1747.329    Cond. No.                      3.11e+07
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 3.11e+07. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [53]: col = ['stock_market_idx']
         df1.drop(col, axis=1, inplace=True)

In [54]: X = np.column_stack((df1['budget'],
                              df1['characters_per_longest_review'], df1['longest_facebook_comm
                              df1['num_critic_for_reviews'], df1['duration'], df1['num_voted_us
                              df1['made_up_column'], df1['num_user_for_reviews']]))

         y = df1['imdb_score']
         X2 = sm.add_constant(X)
         est = sm.OLS(y, X2)
         est2 = est.fit()
         print(est2.summary())

```

#### OLS Regression Results

```

=====
Dep. Variable:          imdb_score    R-squared:                0.948
Model:                  OLS          Adj. R-squared:            0.948
Method:                 Least Squares    F-statistic:              1.154e+04
Date:                  Sat, 20 Jul 2019    Prob (F-statistic):       0.00
Time:                  17:40:01          Log-Likelihood:           8184.5
No. Observations:      5043             AIC:                    -1.635e+04
Df Residuals:          5034             BIC:                    -1.629e+04
Df Model:               8
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.9687	0.091	10.589	0.000	0.789	1.148
x1	-0.0033	0.000	-6.688	0.000	-0.004	-0.002
x2	-0.0136	0.013	-1.049	0.294	-0.039	0.012
x3	-0.0008	0.001	-1.076	0.282	-0.002	0.001
x4	-0.0059	0.001	-5.975	0.000	-0.008	-0.004
x5	0.0068	0.003	2.284	0.022	0.001	0.013
x6	0.0041	0.001	5.605	0.000	0.003	0.006
x7	1.5940	0.006	280.140	0.000	1.583	1.605
x8	0.0037	0.001	3.668	0.000	0.002	0.006

Omnibus:	13728.449	Durbin-Watson:	1.985
Prob(Omnibus):	0.000	Jarque-Bera (JB):	583620066.860
Skew:	33.250	Prob(JB):	0.00
Kurtosis:	1668.253	Cond. No.	3.18e+03

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The condition number is large, 3.18e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [55]: col = ['characters_per_longest_review']
         df1.drop(col, axis=1, inplace=True)
```

```
In [56]: X = np.column_stack((df1['budget'],
                             df1['longest_facebook_comment_review_char'],
                             df1['num_critic_for_reviews'], df1['duration'], df1['num_voted_u
                             df1['made_up_column'], df1['num_user_for_reviews']]))

y = df1['imdb_score']
X2 = sm.add_constant(X)
est = sm.OLS(y, X2)
est2 = est.fit()
print(est2.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          imdb_score    R-squared:                0.948
Model:                  OLS          Adj. R-squared:            0.948
Method:                 Least Squares    F-statistic:            1.319e+04
Date:                  Sat, 20 Jul 2019    Prob (F-statistic):      0.00
Time:                  17:40:02          Log-Likelihood:          8183.9
No. Observations:      5043             AIC:                   -1.635e+04
Df Residuals:          5035             BIC:                   -1.630e+04
Df Model:               7
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.8740	0.015	59.630	0.000	0.845	0.903
x1	-0.0032	0.000	-6.657	0.000	-0.004	-0.002
x2	-0.0008	0.001	-1.110	0.267	-0.002	0.001
x3	-0.0058	0.001	-5.962	0.000	-0.008	-0.004
x4	0.0068	0.003	2.268	0.023	0.001	0.013
x5	0.0041	0.001	5.598	0.000	0.003	0.006
x6	1.5938	0.006	280.199	0.000	1.583	1.605
x7	0.0037	0.001	3.666	0.000	0.002	0.006

=====

Omnibus:	13731.611	Durbin-Watson:	1.985
Prob(Omnibus):	0.000	Jarque-Bera (JB):	584547592.701
Skew:	33.270	Prob(JB):	0.00
Kurtosis:	1669.576	Cond. No.	489.

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [57]: col = ['longest_facebook_comment_review_char']
         df1.drop(col, axis=1, inplace=True)
```

```
In [58]: X = np.column_stack((df1['budget'], df1['num_critic_for_reviews'], df1['duration'], df1['made_up_column'], df1['num_user_for_reviews']))
         y = df1['imdb_score']
         X2 = sm.add_constant(X)
         est = sm.OLS(y, X2)
         est2 = est.fit()
         print(est2.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          imdb_score    R-squared:                0.948
Model:                  OLS          Adj. R-squared:            0.948
Method:                 Least Squares   F-statistic:              1.539e+04
Date:                  Sat, 20 Jul 2019   Prob (F-statistic):       0.00
Time:                  17:40:04          Log-Likelihood:           8183.3
No. Observations:      5043             AIC:                    -1.635e+04
Df Residuals:          5036             BIC:                    -1.631e+04
Df Model:               6
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.8689	0.014	62.356	0.000	0.842	0.896
x1	-0.0032	0.000	-6.660	0.000	-0.004	-0.002
x2	-0.0058	0.001	-5.964	0.000	-0.008	-0.004
x3	0.0068	0.003	2.263	0.024	0.001	0.013
x4	0.0042	0.001	5.614	0.000	0.003	0.006
x5	1.5937	0.006	280.208	0.000	1.583	1.605
x6	0.0036	0.001	3.643	0.000	0.002	0.006

=====

Omnibus:	13733.185	Durbin-Watson:	1.985
Prob(Omnibus):	0.000	Jarque-Bera (JB):	585152979.355
Skew:	33.280	Prob(JB):	0.00
Kurtosis:	1670.440	Cond. No.	445.

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Note: The target variable or dependent variable and independent variables were continuous. I used log transformation to normalize the variables. After verifying Linear Regression assumptions, I move forward with predictive modeling and got adjusted R-square of 95%. After using backward selection technique, the remaining independent variables p-value is < 0.05 which shows that variables are significant. RMSE is also less which is 0.04.

After checking RMSE of the other algorithm, Random Forest root mean squared error is less which is 0.007. It shows that random forest perform better comparing with other models.

<h2> Modules</h2>

```
In [59]: from sklearn import preprocessing
        from sklearn.metrics import precision_recall_fscore_support
        from sklearn.metrics import accuracy_score
        from sklearn.metrics import confusion_matrix
```

<h2>Creating new dataframe as df2</h2>

```
In [60]: df2 = df
```

<h2> Clean & Transform</h2>

The target variable is numeric and have duplicates. My goal is to build a model, which can help

```
In [61]: ### data clean up/ transformation
        # 1 = Good and 0 = when Bad
        df2.loc[df.imdb_score < 8.0, 'imdb_score'] = 0
        df2.loc[df.imdb_score >= 8.0, 'imdb_score'] = 1
```

```
In [62]: df2.imdb_score.unique()
```

```
Out[62]: array([0., 1.])
```

Now, imdb\_score is dichotomous now.

```
In [63]: #create dummy variables
        df2 = pd.get_dummies(df2, columns = ['color', 'director_name',
        'actor_2_name', 'genres', 'actor_1_name', 'movie_title', 'actor_3_name',
        'plot_keywords', 'language', 'country', 'content_rating'], drop_first = True)
```

Some of the columns have missing values. We can deal with this in a few different ways. The simplest solution is to remove them, though we lose many examples in doing so. Alternatively, we could impute the values, replacing the NaN values with an average (mean or median). For the purpose of this simple notebook, we will simply remove them.

```
In [64]: #drop columns
        df2 = df2.drop(columns=['cast_total_facebook_likes', 'made_up_column', 'priority', 'w
        'movie_imdb_link']).drop_duplicates()
```

```
In [65]: def calculate_metrics(y_true,y_pred):
    print(precision_recall_fscore_support(y_true, y_pred,average='macro'))
    print(accuracy_score(y_true, y_pred))
    print(confusion_matrix(y_true, y_pred,labels=[1,0]))

    #df = pd.read_excel("../data/dataset_exercise.xlsx",header=0)
    #df = df.drop("id",axis=1)
    target = df2["imdb_score"]
    df2_x = df2.drop("imdb_score",axis=1)

    df2_x = df2_x.dropna(thresh=int(len(df2_x)*0.5), axis=1)
    print(df2_x.shape)
    df2_x = df2_x.fillna(df2_x.mean())

    x = df2_x.values #returns a numpy array
    min_max_scaler = preprocessing.MinMaxScaler()
    # stand = preprocessing.StandardScaler()
    x_scaled = min_max_scaler.fit_transform(x)
    df2_x_pre = pd.DataFrame(x_scaled)

(5043, 21780)
```

```
In [66]: from sklearn.feature_selection import SelectKBest
    from sklearn.feature_selection import chi2
    bestfeatures = SelectKBest(score_func=chi2, k=10)
    fit = bestfeatures.fit(df2_x_pre,target)
    df2scores = pd.DataFrame(fit.scores_)
    df2columns = pd.DataFrame(df2_x_pre.columns)
    #concat two dataframes for better visualization
    featureScores = pd.concat([df2columns,df2scores],axis=1)
    df2_x_backup = df2_x_pre.copy()
```

```
In [67]: df2_x_pre = df2_x_backup.copy()
    # df_new =
    conf_sum = 0
    index = 0
    for i in df2scores.values:
        if i[0]<0.05:
            conf_sum+=i[0]
            df2_x_pre = df2_x_pre.drop(df2_x_pre.columns[index], axis=1)
            index-=1
        index+=1

    df2_x_pre.shape
```

```
Out [67]: (5043, 21745)
```

<h2>2. Split Data on df2 dataset</h2>

```
In [68]: x_train, x_test, y_train, y_test = train_test_split(df2_x_pre, target, test_size=0.20)

print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)

print("testing data=")
print("1=",np.sum(y_test))
print("0=",len(y_test)-np.sum(y_test))

(4034, 21745)
(1009, 21745)
(4034,)
(1009,)
testing data=
1= 72.0
0= 937.0
```

## <h2> Logistic Regression</h2> (Parametric Model)

```
In [69]: from sklearn.linear_model import LogisticRegression
clf1 = LogisticRegression(random_state=0, solver='lbfgs').fit(x_train, y_train)
y_pred= clf1.predict(x_test)

calculate_metrics(y_test,y_pred)

(0.8927698032961191, 0.6095102573224238, 0.6605154057151981, None)
0.9415262636273538
[[ 16  56]
 [  3 934]]
```

## <h2> Performance Matrix of Logistic Regression </h2>

```
In [70]: from sklearn.metrics import classification_report
predictions = clf1.predict(x_test)
print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0.0	0.94	1.00	0.97	937
1.0	0.84	0.22	0.35	72
avg / total	0.94	0.94	0.93	1009



```
In [71]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,predictions)
```

```
Out[71]: 0.9415262636273538
```

<h2> K Nearest Neighbor</h2> (Nonparametric Model)

```
In [72]: from sklearn.neighbors import KNeighborsClassifier
clf2 = KNeighborsClassifier(n_neighbors=3).fit(x_train, y_train)
y_pred= clf2.predict(x_test)
calculate_metrics(y_test,y_pred)
```

```
(0.7628676470588236, 0.5657091189375074, 0.594474636098345, None)
0.931615460852329
[[ 10  62]
 [  7 930]]
```

<h2> Performance Matrix of KNN </h2>

```
In [73]: predict2 = clf2.predict(x_test)
print(classification_report(y_test,predict2))
```

	precision	recall	f1-score	support
0.0	0.94	0.99	0.96	937
1.0	0.59	0.14	0.22	72
avg / total	0.91	0.93	0.91	1009

```
In [74]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,predict2)
```

```
Out[74]: 0.931615460852329
```

<h2> Gaussian Process Classifier</h2> (Nonparametric Model)

```
In [75]: from sklearn.gaussian_process import GaussianProcessClassifier
clf3 = GaussianProcessClassifier(random_state=0).fit(x_train, y_train)
y_pred= clf3.predict(x_test)
calculate_metrics(y_test,y_pred)
```

```
(0.9652432969215492, 0.5138888888888888, 0.5090229118006895, None)
0.9306243805748265
[[  2  70]
 [  0 937]]
```

## <h2> Performance Matrix of Gaussian </h2>

```
In [76]: predict3 = clf3.predict(x_test)
         print(classification_report(y_test,predict3))
```

	precision	recall	f1-score	support
0.0	0.93	1.00	0.96	937
1.0	1.00	0.03	0.05	72
avg / total	0.94	0.93	0.90	1009

```
In [77]: from sklearn.metrics import accuracy_score
         accuracy_score(y_test,predict3)
```

```
Out[77]: 0.9306243805748265
```

## <h2> Multinomial Naive Bayes</h2> (Parametric Model)

```
In [78]: from sklearn.naive_bayes import MultinomialNB
         clf4 = MultinomialNB().fit(x_train, y_train)
         y_pred= clf4.predict(x_test)
         calculate_metrics(y_test,y_pred)
```

```
(0.4643211100099108, 0.5, 0.4815005138746146, None)
0.9286422200198216
[[ 0 72]
 [ 0 937]]
```

## <h2> Performance Matrix of MNB </h2>

```
In [79]: predict4 = clf4.predict(x_test)
         print(classification_report(y_test,predict4))
```

	precision	recall	f1-score	support
0.0	0.93	1.00	0.96	937
1.0	0.00	0.00	0.00	72
avg / total	0.86	0.93	0.89	1009

```
In [80]: from sklearn.metrics import accuracy_score
         accuracy_score(y_test,predict4)
```

```
Out[80]: 0.9286422200198216
```

## <h2> Decision Tree</h2> (Nonparametric Model)

```
In [81]: from sklearn import tree
         clf5 = tree.DecisionTreeClassifier().fit(x_train, y_train)
         y_pred= clf5.predict(x_test)
         calculate_metrics(y_test,y_pred)
```

```
(0.7666868015705225, 0.7355923159018143, 0.7499380421313506, None)
0.9375619425173439
[[ 36  36]
 [ 27 910]]
```

## <h2> Performance Matrix of Decision Tree </h2>

```
In [82]: predict5 = clf5.predict(x_test)
         print(classification_report(y_test,predict5))
```

	precision	recall	f1-score	support
0.0	0.96	0.97	0.97	937
1.0	0.57	0.50	0.53	72
avg / total	0.93	0.94	0.94	1009

```
In [83]: from sklearn.metrics import accuracy_score
         accuracy_score(y_test,predict5)
```

```
Out[83]: 0.9375619425173439
```

## <h2> Random\_Forest</h2> (Nonparametric Model)

```
In [84]: from sklearn.ensemble import RandomForestClassifier
         clf6 = RandomForestClassifier(n_estimators=10, max_depth=None,min_samples_split=2, ra
         y_pred= clf6.predict(x_test)
         calculate_metrics(y_test,y_pred)
```

```
(0.9703815261044177, 0.5902777777777778, 0.6376799245305985, None)
0.9415262636273538
[[ 13  59]
 [  0 937]]
```

## <h2> Performance Matrix of Random Forest </h2>

```
In [85]: predict6 = clf6.predict(x_test)
         print(classification_report(y_test,predict6))
```

	precision	recall	f1-score	support
0.0	0.94	1.00	0.97	937
1.0	1.00	0.18	0.31	72
avg / total	0.94	0.94	0.92	1009

```
In [86]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,predict6)
```

```
Out[86]: 0.9415262636273538
```

<h2> Gradient Boosting</h2> (Nonparametric Model)

```
In [87]: from sklearn.ensemble import GradientBoostingClassifier
clf7 = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1, r
y_pred= clf7.predict(x_test)
calculate_metrics(y_test,y_pred)
```

```
(0.6842335224688165, 0.728129076248073, 0.7030046467018339, None)
0.9117938553022795
[[ 37  35]
 [ 54 883]]
```

<h2> Performance Matrix of GB </h2>

```
In [88]: predict7 = clf7.predict(x_test)
print(classification_report(y_test,predict7))
```

	precision	recall	f1-score	support
0.0	0.96	0.94	0.95	937
1.0	0.41	0.51	0.45	72
avg / total	0.92	0.91	0.92	1009

```
In [89]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,predict7)
```

```
Out[89]: 0.9117938553022795
```

Note: After seeing confusion matrix of parametric model, I found that Logistic Regression performed better where accuray is 94% and F1-score is 0.93.

After seeing confusion matrix of nonparametric model, I found that Random Forest performed better where accuracy was 94% and F1-score is 0.92.

## <h2>Recommendation</h2>

### Findings

1. Based on continuous dependent and independent variable, random forest (Regressor) performed better.
2. After leveling the continuous dependent variable and creating dummies of independent variables, random forest performed better.